

Different finite objectives select different evidence

22 world classes; 3 collected interactions; all depth-4 targets audited.


Both native objectives select policies using only depth-3 targets.

- ◆

Native mean (128 targets)
- ▼

Reward-matched control
- ✱

Native maximum (128 targets)



Matched control minimizes M_3 with the reference's Brier return floor.

Native mean vs matched control: 6 wins / 1 tie / 15 losses.

Post-design; saved-policy envelopes and numerical bounds, not all-optima intervals.